

Hammad Ahmad

AI/ML Researcher | Trustworthy NLP · Retrieval-Augmented Generation · Knowledge Graphs

📞 +44 7555 856 155 | ✉️ hammadahmad.ml@gmail.com | 📍 Bradford, UK
[LinkedIn](#) | [GitHub](#) | [Portfolio](#) | [ORCID](#)

RESEARCH PROFILE

AI/ML researcher with an MSc (Merit) in Applied AI and a first-author peer-reviewed publication (Springer). My work centres on reliable, transparent language systems: retrieval-augmented generation grounded in knowledge graphs to reduce hallucination, evaluated against benchmarks I build and release, alongside predictive modelling on high-dimensional real-world health data. I am seeking a PhD to develop this into rigorous research on trustworthy, knowledge-grounded NLP.

RESEARCH INTERESTS

Knowledge-grounded and retrieval-augmented NLP; hallucination reduction through structured grounding; evaluation frameworks for RAG systems; interpretable machine learning on high-dimensional health data.

EDUCATION

University of Bradford

Sep 2024 - Sep 2025

MSc, Applied Artificial Intelligence and Data Analytics (Merit)

Bradford, UK

- **Dissertation:** FinLaw-UK: A Graph-Augmented Retrieval Chatbot for Reliable and Transparent UK Financial Regulation. Supervisors: Dr Tillal Eldabi and Dr Irfan Mehmood. Detailed under Research Experience below.
- Selected modules: Artificial Intelligence and Data Science (79), Business Data Analytics (79), Responsible AI: Ethics, Law and Governance (75).

COMSATS University Islamabad

Sep 2020 - Jul 2024

BS, Bioinformatics

Islamabad, Pakistan

- **Thesis:** AI-Assisted Analysis and Prediction of At-Risk Diabetic Individuals (final year research project graded A). Supervisors: Dr Maleeha Azam and Dr Muhammad Umar Khan.

PUBLICATIONS AND RESEARCH OUTPUTS

Ahmad, H., Khan, M.U., & Azam, M. (2024). Comparative Analysis of Machine Learning Methods for Enhancing Sleep Efficiency and Prediction. *ICSMAI 2024* (Springer). DOI: [10.1007/978-3-031-66854-8_1](#)

Software and data. FinLaw-UK implementation and a 110-item UK financial-regulation QA benchmark (80 factual questions, 20 document tasks, 10 case scenarios), publicly released with a measurement-integrity report documenting which figures reproduce and which do not: [FinLaw-UK](#)

RESEARCH EXPERIENCE

University of Bradford

Jan 2025 - Sep 2025

MSc Research Project, Machine Learning & LLMs

Bradford, UK

- Designed and evaluated **FinLaw-UK**, a hybrid architecture combining retrieval-augmented generation, a Neo4j knowledge graph and a locally served Mistral 7B to reduce hallucination in UK financial-regulation question answering, grounded in the FCA Handbook, PRA Rulebook, FRC standards and statutory instruments. Retrieval pipeline: clause-level segmentation of regulatory text, Sentence Transformer embeddings, cross-encoder re-ranking and graph-based citation verification.
- Constructed a 110-item evaluation benchmark (80 factual questions, 20 document tasks, 10 case scenarios) and assessed the pipeline with RAGAS alongside five custom lexical metrics, behind a strict graph-based verification gate that flagged and refused fabricated citations rather than answering.
- Re-measured the submitted evaluation post-hoc and found the reported source-accuracy and citation-quality figures were regex shape-checks, not correctness measures: the scorer awarded a flat 0.85 to any citation-shaped string, with 103 of 110 rows carrying that constant. Published the correction, the true graph-verified citation rate of 3 in 110, and a measurement-integrity report stating which figures reproduce and which do not.
- Quantified the cost of refusing: headline answer relevancy falls 23 points because RAGAS scores a refusal at zero by construction, but excluding the 30 refusals the mean exceeds the non-refusing baseline (0.658 against 0.641), which scored higher only by confabulating answers to unanswerable questions.
- Developed reproducible experimental pipelines with structured evaluation protocols, extending RAGAS with a custom legal-completeness metric that reproduces at 0.68 across both evaluation tracks.

COMSATS University Islamabad

Research Intern, Data Science

Jul 2023 - Jul 2024

Islamabad, Pakistan

- Benchmarked 11 classifiers for diabetes risk prediction on a 253,680-record CDC BRFSS dataset (14% positive class), comparing random over-sampling, SMOTE and ADASYN for class imbalance with resampling confined to the training folds; Random Forest performed best on ROC-AUC and sensitivity.
- Analysed 20+ demographic, lifestyle and clinical indicators, identifying age, general health, BMI, blood pressure and income as the strongest diabetes correlates, and deployed the model as a public risk-screening web app (React, Flask): [DiabetesSense](#).
- Led a first-author comparative study of sleep-efficiency prediction across four regression models (best model Random Forest, R^2 0.86), published at [Springer](#).

INDUSTRY EXPERIENCE

Jobzyl (self-employed)

May 2026 - Present

Founder & Sole Engineer, [jobzyl.com](#)

Remote

- Designed, built and operate a live job-aggregation platform serving a 617K-posting index: cache-first reads with live parallel fan-out to 20 third-party sources on miss, per-source failure isolation, server-side deduplication and streamed incremental results; privacy-preserving client-side CV parsing with encrypted storage and hard-delete data flows. Next.js, TypeScript, Python, FastAPI.

Quality Control Assistant, Myton Food Group (Morrisons), Bradford, UK

Apr 2026 - Present

Outlyst

Oct 2025 - Mar 2026

AI / Machine Learning Engineer

Leeds, UK (Remote)

- Developed backend dialogue-flow logic and workflow automation for an AI voice-agent system (Retell AI, FastAPI), integrating LLM-based dialogue with production infrastructure across 2,100+ calls.
- Reduced average end-to-end response latency 54%, from 2.4s to 1.1s, through async I/O profiling and connection pooling, a decisive factor in conversational systems where response delay degrades interaction quality.

TECHNICAL SKILLS

Machine Learning & NLP : LLMs, Retrieval-Augmented Generation (RAG), Sentence Transformers, cross-encoder re-ranking, RAGAS evaluation, predictive modelling, ensemble methods, PyTorch, scikit-learn

Knowledge & Data : Neo4j knowledge graphs, vector embeddings, semantic search, PostgreSQL, pandas, NumPy

Programming & Tools : Python, SQL, TypeScript, FastAPI, Docker, Git, Linux, Ollama

LANGUAGES

English (IELTS 7.0, CEFR C1) | Urdu (Native) | German (A1.2)

REFEREES

Dr Irfan Mehmood, Associate Professor in Business Analytics and Artificial Intelligence, University of Bradford

i.mehmood@bradford.ac.uk | +44 (0)1274 232646

Professor Tillal Eldabi, Professor of Business Analytics, University of Bradford | t.a.eldabi@bradford.ac.uk (further referees available on request)